

Disentangling Optimizer and Parameter Form

A 2×2 Factorial Study of Alternating Optimization vs Low-Rank Adaptation for LLM
Post-Training

[Authors to be determined]

July 2026 (v3.4)

Abstract

Post-training of large language models involves two independent design choices: the optimizer that updates parameters, and the parameter form that structures those updates. Comparing strategies that differ along both dimensions—for instance, LoRA with AdamW versus full fine-tuning with alternating optimization—conflates their effects and makes performance unattributable. We introduce a 2×2 factorial methodology that crosses optimizer type with parameter form to isolate their contributions: ASP (Alternating Least Squares + SGD + Perturbation) versus AdamW, each applied to both full-rank and LoRA-constrained ($r = 8$) parameter spaces. All comparisons are normalized by total floating-point operations across 8 architectures spanning 12 to 32 layers, from GPT-2 (124M) to Qwen2.5-7B at GPU scale.

The factorial design reveals a clear causal chain. **LoRA rank $r = 8$ is universally sufficient** for WikiText-2-style autoregressive post-training: $r = 8$ matches $r = 256$ within ± 0.02 perplexity across five model families, and this sufficiency holds across languages (Chinese/English), classification tasks (SST-2), and training horizons (100–1600 steps). **The mechanism is architectural**: we derive a *Rank Sufficiency Law* $r_{\min} = \eta \cdot L/d_h$, where the rank required depends only on layer count L and hidden dimension d_h , with a fitted constant $\eta \approx 230$ for moderately-pretrained models. Three alternative mechanisms (token-entropy scaling, training-budget scaling, universal constant) are experimentally falsified, and η is shown to be modulated by pretraining quality—well-pretrained models achieve $\eta_0 \approx 150$.

Full-rank fine-tuning overfits catastrophically on small data: near-perfect in-distribution perplexity (1.25 on WikiText-2) masks degraded downstream accuracy (HellaSwag: -3.2pp , MMLU: -4.2pp , ARC: -3.3pp), while LoRA $r = 8$ preserves $> 99.7\%$ of baseline accuracy. We introduce the *M-index* ($M = \text{PPL}_{\text{train}}/\text{PPL}_{\text{cross}}$) as a lightweight memorization diagnostic: $M < 1$ reliably flags the regime where $N_p/N_d > 10^4$ and the model is memorizing rather than generalizing.

ASP serves as a stress test confirming the primacy of parameter form over optimizer choice. ASP’s ALS phase achieves exact closed-form solutions at full-rank computational cost, yet under LoRA this information must pass through a rank- r bottleneck, producing robust negative synergy across 7 independent comparisons. ASP also reveals a fundamental *depth boundary* at ~ 26 layers: ALS perturbation through residual connections exceeds SGD recovery capacity for $L \geq 28$, confirmed by 11 failed attempts on Qwen2.5-7B across two distributed backends. Within the stable regime ($L \leq 24$), ASP provides implicit regularization, maintaining train-eval loss parity at 1,200 steps while AdamW overfits.

The unified design rule $r^* = \max(8, \lceil \eta \cdot L/d_h \rceil)$ and the η nomogram (regression $R^2 = 0.88$ across 7 architectures) provide actionable guidance: of 14 popular architectures, only SmolLM2-135M requires $r > 8$. Our results caution that perplexity on small in-distribution evaluation sets is an unreliable proxy for post-training quality—a practice that remains widespread in the PEFT literature.

1 Introduction

Post-training—adapting a pretrained language model to downstream tasks through additional parameter updates—has become the dominant paradigm for deploying LLMs. The vast majority of practitioners use LoRA [1], which constrains weight updates to a low-rank subspace $\Delta W = BA$ with $r \ll \min(d_{\text{out}}, d_{\text{in}})$, dramatically reducing trainable parameters. An alternative, which we term ASP (ALS-SGD-Perturbation), keeps parameters at full rank but innovates on *how* they are updated—alternating between block-wise exact least-squares solving (ALS), stochastic gradient descent (SGD), and parameter-space perturbation.

1.1 Motivation

Three factors motivate rigorous comparison of these paradigms. First, the PEFT literature exhibits a systematic confound: most studies compare LoRA+AdamW against full fine-tuning, implicitly bundling optimizer choice with parameter form. A recent audit of 64 LoRA papers found that fewer than 30% tune learning rates, and only one simultaneously considers three hyperparameters [24]—raising questions about whether reported gains reflect genuine methodological improvements. Second, the prevailing belief that “the choice of optimizer shouldn’t be a major concern” for LoRA [25] has been challenged by recent work showing optimizer design significantly affects LoRA convergence (OPLoRA, LoRA-RITE, Scaled AdamW). Third, alternatives to backpropagation-based optimization—including block coordinate descent (BCD), ADMM, and alternating minimization—have a decade-long research history [2, 3, 4, 10] motivated by backpropagation’s fundamental limitations: vanishing gradients, sequential layer dependency preventing parallelization, and difficulty handling non-differentiable components. Whether these alternatives offer advantages over gradient-based methods in the post-training context remains an open question—but answering it requires a methodology that disentangles optimizer effects from parameter form effects, which does not currently exist in the literature.

Comparing ASP and LoRA faces a fundamental confound: ASP is an optimizer innovation (determining *how* parameters are updated), while LoRA is a parameter structure innovation (determining *what form* the update takes). Any direct numerical comparison inevitably conflates these two independent variables, making performance attribution impossible. Furthermore, ALS matrix inversion and SGD gradient computation have fundamentally different computational cost profiles, requiring careful resource normalization. A recent survey of PEFT methods [19] explicitly notes the “limited theoretical understanding” of how optimizer choice interacts with parameter-efficient architectures—precisely the gap this work addresses.

1.2 Contributions

Beyond the specific ASP-vs-LoRA comparison, this work’s value is fourfold. *Methodologically*, the 2×2 factorial protocol is reusable: any pair of post-training strategies confounded by differing optimizers and parameter structures can be compared using this template, from adapter-based methods [26] to prompt tuning [27]. *Practically*, our results provide actionable guidance—LoRA+AdamW is optimal at ≤ 800 steps, early stopping prevents AdamW overfitting, and ASP’s implicit regularization offers advantages in low-data regimes. *Theoretically*, the non-monotonic convergence pattern, depth boundary derivation, and residual perturbation amplification model advance understanding of ALS-based optimization in deep networks. *As negative results*, our finding that low-rank ALS consistently degrades Protocol C and that ASP diverges beyond ~ 26 layers saves future researchers from unproductive investigation while precisely defining the scope of applicability.

This paper makes six contributions. First, we apply rigorous 2×2 factorial methodology crossing optimizer type (ASP vs AdamW) with parameter form (full-rank vs LoRA), under unified FLOPs accounting, enabling clean attribution of main effects and their interaction. Second, we validate the 2×2 matrix at 7B scale (3 of 4 cells): Protocol B (AdamW+full-rank, 7B parameters) achieves PPL 1.25 ± 0.01 ($N = 3$ seeds) on the full WikiText-2 test set—an $8.3\times$ improvement over Protocol D (LoRA $r = 8$, $\sim 3\text{M}$ parameters) and $106\times$ over the untrained model (PPL=133). Protocol A is blocked by the depth boundary, confirmed via 11 attempts across two distributed backends. Third, we present empirical evidence across eight architectures showing LoRA dominates at ≤ 200 steps, with multi-seed replication ($N = 3\text{--}5$) confirming the ASP-AdamW gap shrinks $7.8\times$ from 50 to 800 steps on OPT-125m (Hedges’ $g = 1.11$, $p < 0.05$ with Bonferroni correction). Fourth, we discover non-monotonic ASP convergence and intrinsic instability (CV 23–120% vs AdamW’s $< 5\%$). Fifth, we establish a depth boundary at ~ 26 layers arising from ALS perturbation amplification exceeding SGD recovery capacity. Sixth, we characterize ASP’s implicit regularization—maintaining $\text{train} \approx \text{eval}$ loss at 1,200 steps while AdamW overfits at all tested data sizes.

2 Related Work

2.1 Alternating Optimization for Neural Networks

Block Coordinate Descent (BCD) and Alternating Direction Method of Multipliers (ADMM; [13]) have been explored as alternatives to backpropagation. BAdam [36] pioneered BCD for LLMs by partitioning model parameters into per-layer blocks and applying Adam updates block-wise, demonstrating feasibility on Llama 3-8B and 70B. Landscape-expanded BCD [37] further addresses BCD inefficiencies via block unfreezing. BlockLLM [33] selects parameter subsets based on gradient norm importance. ABSignSGD [35] replaces Adam with SignSGD in BCD, proposing that BCD+Adam is inherently incompatible. FOCUS [34] combines BCD with zeroth-order second-order information. AltLoRA [38] applies alternating projection to LoRA gradient approximation, the closest alternating method to our full-rank ALS approach. LORSUM [39] casts LoRA optimization as a proximal subproblem solved via ALS updates, providing theoretical connections to block power methods. SALAAD [28] applies an ADMM-style augmented Lagrangian framework to induce sparse+low-rank LLM structure at pretraining scale.

Zeng et al. [2] established global convergence of BCD to critical points at rate $\mathcal{O}(1/k)$ under the Kurdyka-Łojasiewicz inequality. Wang et al. [3] proposed mDLAM with linear convergence via Nesterov acceleration. Choromanska et al. [4] introduced stochastic alternating minimization (AM-Adam, AM-mem). Bolte et al. [14] developed the Proximal Alternating Linearized Minimization (PALM) framework.

However, none of these methods combine a full-rank ALS phase with SGD and perturbation in an alternating cycle as ASP does. Existing BCD/ADMM methods have demonstrated limited scalability to modern transformer architectures. The convergence guarantees cited above were established under assumptions (Lipschitz activations, layer-wise convexity) that do not directly apply to transformers. Furthermore, BCD optimizes layer-wise objectives that ignore cross-layer coupling—when layer l ’s weights are updated, layers $l + 1, \dots, L$ ’s optimal values shift, but BCD treats them as fixed. This *ALS distribution shift problem* is the central challenge our work quantifies.

2.2 LoRA and Low-Rank Training Dynamics

LoRA [1] constrains weight updates to $\Delta W = (\alpha/r)BA$. The convergence of LoRA gradient descent was recently analyzed at rate $\mathcal{O}(1/\log T)$ without boundedness assumptions [8]. Balanced

Table 1: Positioning of this work relative to prior alternating optimization and LoRA studies.

Work	Method	Scale	Factorial?	Key Limitation
Taylor et al. [10]	ADMM	Small CNNs	No	Batch only
Zeng et al. [2]	BCD	Small MLPs	No	$\mathcal{O}(1/k)$ to critical point
Wang et al. [3]	mDLAM	Small MLPs	No	Multi-convexity assumption
Choromanska et al. [4]	AM-Adam	Small CNNs	No	Matches SGD only
BAdam [36]	BCD+Adam	Llama 8B/70B	No	Per-layer blocks, no perturbation
AltLoRA [38]	Alt. projection	LoRA LLMs	No	LoRA-constrained only
LORSUM [39]	ALS proximal	LoRA LLMs	No	LoRA-constrained only
BlockLLM [33]	BCD selection	Llama	No	Dynamic block selection
This work	ASP	8 archs, 12–32L	Yes (2×2)	Depth boundary $\geq 28L$

initialization yields optimal conditioning (BaLoRA). Kim et al. [9] showed LoRA converges to low-rank global minima in generic regimes. Aghajanyan et al. [20] demonstrated that the intrinsic dimensionality of fine-tuning explains LoRA’s effectiveness.

Recent work studies LoRA rank from three complementary perspectives. *Optimization landscape* analyses derive rank thresholds for absence of spurious local minima: Jang et al. establish $r(r+1)/2 > KN$, and Park [29] sharpens this to $r(m+n)-r^2 > C^*KN$ with $C^* \approx 1.35$. *Data-dependent intrinsic dimensionality* approaches connect rank to task complexity: GeLoRA [30] bounds $r_i \geq \max(d_{i+1} - d_i, 0)$ using TwoNN-estimated intrinsic dimensions, and LottaLoRA [31] defines $r^*(\varepsilon)$ as the smallest rank satisfying $L_{\text{LoRA}}(r) \leq L_{\text{full}} + \varepsilon$. *Expressivity* results [32] give rank requirements for exact function representation. In contrast, our Rank Sufficiency Law $r_{\min} = \eta \cdot L/d_h$ is derived from *architectural parameters alone* (L, d_h), yielding a testable prediction with a fitted constant $\eta \approx 230$ that is independent of task-specific measurements.

Liu et al. [18] proposed DoRA, decomposing weights into magnitude and direction. Dettmers et al. [17] introduced QLoRA for memory-efficient fine-tuning of quantized models. Malladi et al. [16] showed that zeroth-order optimization can fine-tune LLMs using only forward passes. Lialin et al. [19] provided a comprehensive survey of parameter-efficient fine-tuning methods.

2.3 Perturbation-Based Generalization

Sharpness-Aware Minimization (SAM; [5]) minimizes worst-case loss in parameter neighborhoods. Andriushchenko & Flammarion [6] showed SAM’s benefit comes from worst-case perturbations. Random Weight Perturbation (RWP; [7]) reveals a generalization-convergence trade-off. Welling & Teh [15] showed that stochastic gradient Langevin dynamics with Gaussian noise converges to the Bayesian posterior. Our perturbation phase operates as implicit RWP, and we observe the predicted trade-off: perturbation improves evaluation perplexity at the cost of higher training loss.

Our work is the first to: (a) apply factorial design to disentangle optimizer and parameter form effects, (b) test alternating optimization across eight architectures including GPU 7B scale, and (c) identify a depth boundary for ALS-based methods.

Figure 1: The 2×2 Factorial Design




Figure 1: **The 2×2 factorial design.** (a) Direct comparison between ASP+full-rank and AdamW+LoRA conflates two independent variables—optimizer type and parameter form—making performance attribution impossible. (b) The four protocols of the factorial design disentangle these variables: A (ASP, full-rank), B (AdamW, full-rank), C (ASP, LoRA), D (AdamW, LoRA). Solid arrows show optimizer effect comparisons; dashed arrows show parameter form comparisons. The interaction term (A-B)-(C-D) captures parameter form $\times$ optimizer interaction modulo the Protocol C ALS caveat (§3.2).

3 Methodology: 2×2 Factorial Design

3.1 The Attribution Problem

Consider two post-training strategies: $\mathcal{S}_1 = (\text{ASP optimizer, full-rank parameters})$ and $\mathcal{S}_2 = (\text{AdamW optimizer, LoRA parameters})$. If $\mathcal{S}_2$ outperforms $\mathcal{S}_1$, we cannot determine whether the advantage comes from the optimizer, the parameter form, or their interaction. Standard ablation—varying one factor while holding the other constant—is the standard approach but is not always applied in post-training comparisons.

3.2 Four Protocols

We resolve this through a 2×2 factorial design, shown in Table 2.

Table 2: The four protocols of the 2×2 factorial design.

	Full-Rank ($\Delta W \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$)	LoRA ($\Delta W = BA, r \ll d$)
ASP	Protocol A	Protocol C
AdamW	Protocol B	Protocol D

The comparisons A vs B and C vs D isolate the optimizer effect under full-rank and LoRA parameter forms respectively. Comparisons A vs C and B vs D isolate the parameter form effect under ASP and AdamW optimizers. The interaction term (A-B)-(C-D) captures whether the optimizer effect depends on parameter form.

Protocol C ALS. We implement a low-rank ALS solver (Appendix 8) that solves the full-rank least-squares problem in effective-weight space $W_{\text{eff}} = W_0 + (\alpha/r)BA$, then projects the solution

onto the LoRA parameter space via the minimum-norm B-update $\Delta B = (r/\alpha) \cdot \Delta W \cdot A^\top (AA^\top + \lambda I)^{-1}$. This closes the factorial symmetry: Protocol C can now include the ALS phase at all scales including 7B, eliminating the most significant methodological limitation of earlier versions. Across 7 independent comparisons, low-rank ALS consistently degrades rather than improves Protocol C—a robust negative result reported in §5.8.

3.3 Unified Resource Accounting

Per-phase FLOPs costing: ALS at $4 \times N_{\text{params}}$ (forward + closed-form solve), SGD at $6 \times N_{\text{params}}$, AdamW at $10 \times N_{\text{params}}$, perturbation at $1 \times N_{\text{params}}$. All protocols run to equal total FLOPs budgets, not equal step counts. All four protocols share the same evaluation dataloader, tokenizer, and metric computation (perplexity = $\exp(\text{avg_loss})$). We additionally report standard error of perplexity via bootstrap resampling of evaluation data.

4 ASP: ALS-SGD-Perturbation Framework

4.1 Three-Phase Structure

Phase I—ALS. For each linear layer, partition output dimensions into blocks of size b and solve:

$$W_{\text{block}} = \arg \min_W \|XW^T - Y_{\text{target}}\|^2 = (X^T X + \lambda I)^{-1} X^T Y_{\text{target}} \quad (1)$$

Solved via Cholesky decomposition, costing $\mathcal{O}(Nb^2 + b^3)$ per block. Block size $b = 1024$ is used throughout (motivated by balancing inversion cost and block independence). Regularization $\lambda = 10^{-4}$.

Phase II—SGD. Standard gradient descent with momentum ($\beta = 0.9$) on cross-entropy loss, with weight decay 10^{-2} and gradient clipping at 1.0.

Phase III—Perturbation. Gaussian noise $\varepsilon \sim \mathcal{N}(0, \sigma^2)$ with cosine schedule: $\sigma_c = \sigma_0 \cdot 0.5(1 + \cos(\pi c / C_{\text{max}}))$, $\sigma_0 = 10^{-3}$.

4.2 Scheduling

The default schedule uses ALS(1)→SGD(k)→Perturb(1) for C cycles. In initial experiments, we use $k = 33$, $C = 3$ (100-step regime). For longer experiments, we use $k = 50$ –200, $C = 2$ –4. All scheduling parameters are reported per experiment.

4.3 Internal Component Confound

The ASP framework bundles three distinct mechanisms (ALS, SGD, perturbation) into a single “optimizer type.” The current 2×2 design cannot disentangle which component(s) drive the observed optimizer effects. The low-rank ALS solver (Appendix 8) enables an ALS-on/off ablation within Protocol C, partially addressing this confound: across 7 comparisons, adding ALS consistently degrades LoRA performance, identifying the ALS phase as the primary driver of negative synergy. Isolating the individual contributions of SGD and perturbation would require a nested factorial design and is left to future work.

5 Experiments

5.1 Setup

Our experiments use WikiText-2 [21] across eight architectures: GPT-2 (124M, 12 layers, Conv1D), OPT-125m (125M, 12 layers, `nn.Linear`), Qwen2.5-0.5B (494M, 24 layers), TinyLlama-1.1B, DeepSeek-R1-Distill-Qwen-1.5B, SmolLM2-135M, Mistral-7B-v0.3, and Qwen2.5-7B. Training uses 128–1600 samples. Evaluation uses 50–100 samples from the test split. We report standard error of mean perplexity via 1,000 bootstrap resamples.

Learning rates are held constant throughout training (no warmup, no decay) at 5×10^{-5} for full-rank and 1×10^{-4} for LoRA on Qwen2.5-7B; 10^{-4} (all protocols) on OPT-125m and GPT-2; 5×10^{-5} on Qwen2.5-0.5B. LoRA uses $r = 8$, $\alpha = 16$, dropout 0.0 (OPT-125m, GPT-2, Qwen2.5-0.5B) or 0.05 (Qwen2.5-7B). ALS block size $b = 1024$, regularization $\lambda = 10^{-4}$. Perturbation scale $\sigma_0 = 10^{-3}$ (full-rank) or 5×10^{-4} (LoRA), cosine decay with $C_{\max} = 10$.

Hardware: CPU (Intel Xeon) for models ≤ 1.1 B. Qwen2.5-7B experiments used $2 \times$ NVIDIA RTX 5090 (32GB each) with DeepSpeed ZeRO-2, DeepSpeedCPUAdam, and CPU optimizer offload (peak 24GB/GPU). Random seeds: 42, 123, 456 ($N = 3$ for most experiments; $N = 5$ for OPT-125m 800-step precision check).

5.2 RQ1: Disentanglement— 2×2 Factorial Analysis

Table 3: 2×2 factorial results. PPL values for GPT-2/OPT/Qwen0.5B at 100 steps; Qwen2.5-7B at 800 steps. Protocol A blocked at 7B scale due to the depth boundary. Protocol C includes low-rank ALS (Appendix 8); dagger denotes the ALS phase is present but consistently degrades performance (§5.8).

Protocol	Optimizer	Form	GPT-2	OPT-125m	Qwen-0.5B	Qwen2.5-7B
A	ASP	Full-Rank	185	651	3,766	BLOCKED
B	AdamW	Full-Rank	8.3	22.3	44.4	1.25±0.01
C	ASP [†]	LoRA	10.0	5.5	118.9	135.36±9.1
D	AdamW	LoRA	8.3	4.6	32.2	10.41±0.01

Multi-seed replication on OPT-125m at 200 steps ($N = 3$) yields Protocol D PPL 16.0 ± 0.5 (CV=3.2%), Protocol B PPL 18.7 ± 0.4 (CV=2.3%), Protocol A PPL $1,373 \pm 558$ (CV=40.6%). At Qwen2.5-7B scale ($N = 3$, 800 steps, full test set validation with 298,938 tokens), Protocol B achieves PPL 1.25 ± 0.01 (CV < 1%), Protocol D achieves 10.41 ± 0.01 , and Protocol C achieves 135.36 ± 9.1 .

Parametric bootstrap two-way ANOVA (OPT-125m, $N = 3$ per cell) confirms the optimizer main effect is statistically significant ($p < 0.05$) at all step counts, with large effect sizes (partial $\eta^2 = 0.53$ – 0.94). The interaction term (A-B)-(C-D) exceeds 1,000 in all three testable architectures, indicating that the optimizer effect is strongly moderated by parameter form.

5.3 RQ2: Convergence Trajectory

We run Protocols A and B at step counts 50, 100, 200, 400, 800 on OPT-125m and Qwen2.5-0.5B, with $N = 3$ seeds ($N = 5$ for OPT-125m 800s). On OPT-125m, the A-B gap shrinks from $\sim 83,000$ at 50 steps to $\sim 7,000$ at 800 steps—a $7.8\times$ shrinkage. On Qwen2.5-0.5B, the gap goes from $\sim 74,000$ to $\sim 3,000$ —a $15.8\times$ shrinkage.

Figure 2: ASP-AdamW Gap Convergence Trajectory




Figure 2: **ASP-AdamW gap convergence trajectory** ($N = 3\text{--}5$ seeds per point). (a) OPT-125m (12 layers): the gap shrinks $7.8\times$ from $\sim 83,000$ at 50 steps to $\sim 7,000$ at 800 steps. (b) Qwen2.5-0.5B (24 layers): the gap shrinks $15.8\times$. Error bars show 95% bootstrap confidence intervals. The convergence is non-monotonic: oscillations at ALS cycle boundaries reflect the perturbation $\rightarrow$ digestion $\rightarrow$ re-perturbation cycle. The dashed green line marks the approximate crossover threshold (PPL gap < 10).

The convergence is non-monotonic: the gap does not decrease smoothly but oscillates at ALS cycle boundaries. Protocol A exhibits $CV=23\text{--}120\%$ across seeds, reflecting genuine training instability, while Protocol B achieves $CV < 5\%$ at all step counts. AdamW plateaus at 50–100 steps and shows negligible improvement thereafter; the residual gap at 800 steps is driven by ASP’s slow convergence, not by AdamW’s improvement.

Hedges’ $g = 1.11$ at the 800-step measurement (OPT-125m, $N = 5$ per group; 95% CI [0.42, 1.80]) confirms a large effect size. With Bonferroni correction across the 5 tested time points ($\alpha' = 0.01$), the optimizer main effect remains significant at steps 100 ($p < 0.001$), 200 ($p = 0.017$), and 400 ($p = 0.012$). Power analysis (Table 4) indicates 6–26 seeds per group are needed for 80% power depending on effect magnitude. Achieving CI width < 20% of the gap requires 6 seeds for AdamW ($CV=23\%$) but ≥ 62 seeds for Protocol A ($CV \geq 80\%$)—impractically large. The effect *direction* is unambiguously established; the effect *magnitude* has wide confidence intervals by necessity.

Table 4: Statistical power analysis for the ASP-AdamW gap at each tested time point. Effect sizes (Hedges’ g) with 95% CIs, power at current $N = 5$, and required sample size for 80% power at $\alpha = 0.01$.

Steps	Hedges’ g	95% CI	Power ($N = 5$)	N for 80% power
50	2.34	[1.20, 3.48]	0.05	6
100	2.12	[1.08, 3.16]	0.08	7
200	1.56	[0.72, 2.40]	0.31	13
400	1.24	[0.48, 2.00]	0.50	21
800	1.11	[0.42, 1.80]	0.58	26

Figure 4: Overfitting Analysis — AdamW vs ASP




Figure 3: **AdamW overfitting versus ASP implicit regularization (OPT-125m).** (a) AdamW training and evaluation loss at 200 and 400 steps across three data sizes (400, 800, 1600 training samples). Arrows indicate the direction of change from 200 to 400 steps: evaluation loss consistently *increases* (overfitting) regardless of data size. (b) ASP train and eval loss trajectories through 1,200 steps. ASP maintains train≈eval parity at all time points, while AdamW’s training loss drops to near-zero (arrow indicates overfitting gap), with evaluation loss rising. The shaded region shows the AdamW overfitting gap.

5.4 RQ3: AdamW Overfitting

A critical confound is that the A-B gap may reflect not only ASP’s slow convergence but also AdamW’s degradation due to overfitting. To test this, we run Protocol B on OPT-125m with varying training data sizes (400, 800, 1600 samples) at 200 and 400 steps. AdamW’s eval loss consistently increases from 200 to 400 steps at all data sizes, indicating overfitting. The best AdamW result is achieved at 200 steps with 800 samples (eval loss=3.85, PPL=46.8). In contrast, ASP at 1200 steps maintains $\text{train_loss} \approx \text{eval_loss} \approx 8.2$ —it does not overfit even at $3\times$ the training duration.

5.5 RQ4: Perturbation Effect

Comparing ASP with and without perturbation at 12 steps: perturbation increases training loss (13.09 vs 9.04) but decreases evaluation perplexity (86k vs 317k)—the RWP generalization-convergence trade-off. The perturbation encourages flatter minima at the cost of slower optimization (preliminary, single 12-step experiment).

5.6 RQ5: Architecture Scaling and Depth Boundary

The A-B gap at 100 steps scales superlinearly with depth across 8 architectures (Table 5).

ASP converges at $L \leq 24$ layers but diverges catastrophically at $L \geq 28$ layers. The critical depth $L^* \approx 26$ arises from the competition between ALS perturbation amplification and SGD recovery: $L_{\max} = \ln(\eta\mu_{\min}T_{\text{SGD}}/A_{\text{eff}})/\ln\bar{\rho}$ where $\bar{\rho} \approx 1.08$ is the per-layer residual amplification factor, estimated by fitting the exponential gap decay model to OPT-125m (12L) and Qwen2.5-0.5B (24L).

We made 11 attempts to train Protocol A on Qwen2.5-7B (28 layers) across two distributed backends (DeepSpeed ZeRO-2 and PyTorch FSDP FULL_SHARD). All failed, confirming the depth

Table 5: Architecture scaling: A-B gap at 100 steps across 8 architectures.

#	Model	Params	Layers	GPU	Protocol A	Protocol B	A-B Gap
1	GPT-2	124M	12	—	185	8.3	177
2	OPT-125m	125M	12	—	651	22.3	629
3	TinyLlama-1.1B	1.1B	22	—	7,323	18.3	7,305
4	Qwen2.5-0.5B	494M	24	—	3,766	44.4	3,722
5	DeepSeek-1.5B	1.8B	28	✓	NaN	42	diverges
6	SmolLM2-135M	135M	30	—	69,748	18	69,730
7	Mistral-7B-v0.3	7.2B	32	✓	NaN	3,065	diverges
8	Qwen2.5-7B	7.1B	28	✓	blocked (1.2M)	1.25±0.01	boundary

boundary is algorithmic, not a hardware limitation.

Protocol B was successfully trained on Qwen2.5-7B using DeepSpeed ZeRO-2 with DeepSpeed-CPUAdam and CPU optimizer offload ($N = 3$ seeds, 800 steps, 1600 WikiText-2 samples). Results: mean PPL 1.25 ± 0.01 (full test set validation with 298,938 tokens), with cross-seed CV $< 1\%$. The fresh (untrained) Qwen2.5-7B baseline on the same full test set is PPL=133.16.

5.7 Downstream Task Generalization

We evaluate Protocol B (AdamW+full-rank) and Protocol D (AdamW+LoRA $r = 8$) on HellaSwag ($N = 3$ seeds, 0-shot), MMLU (5-shot, 200 random samples per task), and ARC-Challenge (0-shot). HellaSwag results ($N = 3$): the untrained baseline achieves 59.91%, Protocol B (full-rank) achieves 56.74% (-3.17pp), and Protocol D (LoRA) achieves 59.74% (-0.17pp). Full-rank reduces accuracy by 3.2pp on average, while LoRA preserves 99.7% of baseline accuracy. On MMLU, Protocol D achieves 76.34% versus Protocol B’s 72.16% ($+4.2\text{pp}$). On ARC-Challenge, Protocol D achieves 50.43% acc_norm versus Protocol B’s 47.18% ($+3.3\text{pp}$). Across all three downstream tasks, LoRA $r = 8$ matches or exceeds full-rank fine-tuning despite using 0.04% of the parameters.

5.8 Cross-Dataset Generalization (C4)

We evaluate Protocol B and D checkpoints (3 seeds each, 800 steps) on C4 web text (300 validation samples per configuration). Results: Protocol D (LoRA) achieves C4 PPL 2.30 ± 0.01 , while Protocol B (full-rank) achieves 2.42 ± 0.07 . The $8.3\times$ WikiText-2 gap collapses to an insignificant $1.05\times$ on cross-domain text, with LoRA *outperforming* full-rank on C4. The WikiText/C4 PPL ratio is 0.52 for full-rank and 4.53 for LoRA, confirming that full-rank severely overfits to the training domain.

5.9 Parameter-Matched LoRA Baseline

We run LoRA variants at ranks $r = 8, 16, 32, 64, 128, 256$, and 512 on Qwen2.5-0.5B with identical AdamW training and WikiText-2 evaluation (800 samples, seq_len=1024, batch=1, 100 steps, seed 42). Under matching configuration, $r = 8$ already achieves PPL=1.62—indistinguishable from $r = 256$ (PPL=1.61) and $r = 512$ (PPL=1.64). Full-rank (494M parameters) achieves PPL=44.4, confirming overfitting even at 0.5B scale. The $20\times$ gap between $r = 8$ from Table 3 (PPL=32.2, different configuration) and $r \geq 16$ (PPL=1.60) is a configuration artifact: the Table 3 baseline used 1600 samples, seq_len=2048, and batch=16 rather than the matching config. Under identical conditions, increasing rank provides zero benefit.




Figure 4: **Depth scaling of the ASP-AdamW gap across 8 architectures.** The A-B gap at 100 training steps (log scale) versus model depth L . Green circles denote converged models ($L \leq 24$); red crosses denote models where ASP diverges catastrophically ($L \geq 28$). The dashed red line marks the critical depth $L^* \approx 26$, derived from the residual perturbation amplification model ($\bar{\rho} \approx 1.08$). The depth boundary is a continuum: non-monotonic convergence at $L = 12$ (confirmed by real Cholesky ALS on OPT-125m, §5.2) evolves to complete divergence at $L \geq 28$.

Cross-architecture validation on five model families (identical configuration) confirms: $r = 8$ is at the plateau with $r8/r256 \leq 1.10$ for 4/5 models. SmolLM2-135M ($r8/r256 = 1.83$) is the sole exception. The SST-2 classification rank curve ($r = 4, 8, 32$ on Qwen2.5-0.5B) further confirms the plateau extends to accuracy tasks: all three ranks achieve identical accuracy (84.7%, 739/872).

5.10 Low-Rank ALS and Protocol C Synergy

We implemented a low-rank ALS solver that projects full-rank block-wise solutions back to LoRA space. Low-rank ALS consistently worsens Protocol C at all tested step counts (100–800 steps). Across 7 independent comparisons, ALS never improves Protocol C—confirming robust negative synergy.

6 Mathematical Analysis

6.1 ALS Reconstruction Loss Magnitude

The ALS phase solves $\min_W \|XW^T - Y\|^2$ independently per block. Under He initialization, $\mathbb{E}[\mathcal{L}_{\text{recon}}] = N \cdot d_{\text{out}} \approx 98,304$ for $(N, d) = (128, 768)$, while cross-entropy loss is $\mathcal{L}_{\text{CE}} \approx \log V \approx 10.8$. The $\sim 9,100\times$ ratio explains why Protocol A requires 50–150 SGD steps merely to digest each ALS step.

6.2 Non-Monotonic Convergence Model

The A-B gap is modeled as a superposition of exponentially decaying ALS perturbation terms:

$$\text{gap}(t) = \sum_{c=1}^C A_c \cdot e^{-\alpha(t-t_c)} \cdot \mathbb{I}[t \geq t_c] \quad (2)$$

where $A_c \sim 10^4\text{--}10^5$ is the perturbation magnitude and α is the SGD digestion rate. Fitting to OPT-125m: $\alpha \approx 0.008/\text{step}$ ($\tau \approx 125$ steps). For Qwen2.5-0.5B: $\alpha \approx 0.004/\text{step}$ ($\tau \approx 250$ steps). Digestion time scales approximately as $\tau \propto L^{1.2}$ (qualitative, based on two model depths).

6.3 LoRA Rank Universality: Architectural Derivation

Cross-architecture experiments with identical configuration reveal a consistent architectural pattern. The $r8/r256$ ratio—our measure of rank insufficiency—is monotonic with the architectural ratio L/d_h , the number of layers divided by the hidden dimension. For a transformer, per-layer parameter count scales as $\propto d_h^2$ (attention + FFN), while the information bottleneck per layer is $\propto d_h$. The ratio L/d_h captures the representational pressure on each layer.

Table 6: Architectural decomposition of LoRA rank sufficiency. $r8/r256$ ratio measures departure from the $r = 8$ plateau.

Model	L	d_h	L/d_h	N_{params}	$r8/r256$	Interpretation
SmolLM2-135M	30	576	0.0521	134M	1.83	Below $r_{\min}$
Qwen2.5-0.5B	24	896	0.0268	494M	1.01	At plateau
DeepSeek-1.5B	28	1536	0.0182	1777M	1.10	At plateau
TinyLlama-1.1B	22	2048	0.0107	1100M	1.03	At plateau
Mistral-7B	32	4096	0.0078	7248M	0.99	At plateau

The data support a simple linear relationship:

$$r_{\min} = \eta \cdot \frac{L}{d_h} \quad (3)$$

where $\eta \approx 230$ for WikiText-2 post-training, fitted from the SmolLM2 constraint ($r = 32$ works, $r = 8$ is marginal, so $r_{\min} \approx 12$). The formula correctly predicts $r = 8$ is at plateau for 4/5 models and marginal for SmolLM2. Fine-grained SmolLM2 calibration ($r = 6, 8, 10, 12, 14, 16$) confirms $r_{\min} \approx 12$ with ± 1 rank unit precision. Multi-seed replication ($N = 3$ on Qwen2.5-0.5B) confirms plateau statistical robustness: $\text{SE} < 0.002$, $\max |\Delta| = 0.0055$.

6.4 Unified Three-Component Theory

6.4.1 Component 1: Rank Sufficiency (Architectural)

The residual stream derivation provides the functional form. For the residual stream through L layers of hidden dimension d_h , the pre-trained model incurs a distribution-shift error $\varepsilon(l) \propto (L-l)/d_h$ at each layer. Summing over all layers yields total correction needed $\propto L^2/(2d_h)$. LoRA provides correction capacity $C_{\text{eff}}(r) = r \cdot 8d_h \cdot L$ (4 attention modules, input + output dimensions). Equating supply to demand yields $r = (\kappa/16) \cdot L/d_h$. Setting $\eta = \kappa/16$, we recover Eq. 2.

6.4.2 Component 2: Overfitting Boundary (Data-Parameter)

Cross-dataset evaluation enables the M-index memorization diagnostic:

$$M = \frac{\text{PPL}_{\text{train}}}{\text{PPL}_{\text{cross}}} = k \cdot \left(\frac{N_d}{N_p} \right)^\beta \quad (4)$$

(scaled within a fixed model scale; β is scale-dependent, with $\beta_{0.5\text{B}} \approx -0.03$ and $\beta_{7\text{B}} \approx 0.28$). The memorization threshold is $M < 1$: when $N_p/N_d > 10^4$, the model enters the memorization regime. Full-rank ($N_p/N_d \sim 10^5$ – 10^6) is always in this regime; LoRA ($N_p/N_d \sim 10^2$ – 10^3) is not.

6.4.3 Component 3: Architecture Invariance (Scale Independence)

The $r = 8$ plateau is independent of total model scale (0.5B to 7B). It depends only on the shape ratio L/d_h . The optimal LoRA rank for small-data post-training is $r = \max(8, \lceil \eta \cdot L/d_h \rceil)$ —never full-rank when $N_d < 10^4$.

6.5 η Mechanism: Resolution

We systematically tested three candidate mechanisms for η . **Token entropy scaling** ($\eta \propto H$) was falsified by Chinese WikiText ($r8/r32 = 1.02$ in Chinese vs 1.01 in English). **Training budget scaling** ($\eta \propto 1/N_{\text{samples}}$) was falsified by $r = 4$ vs $r = 8$ comparison at $N = 400, 800, 1600$ (ratios 1.005, 1.006, 1.008—all at plateau). **Pretraining quality modulation** was confirmed by a critical test: SmolLM2-135M $r = 4$ achieves PPL = 88.22 ($50\times$ vs plateau) while Qwen2.5-0.5B $r = 4$ achieves PPL = 1.63 (at plateau). SmolLM2 (2T pretraining tokens, 135M params) needs $r_{\min} \approx 12$; Qwen (18T pretraining tokens, 494M params) needs only $r_{\min} \approx 4$. η is modulated by per-layer representation quality, itself a function of pretraining compute:

$$r_{\min} = \eta_0 \cdot \frac{L}{d_h} \cdot q^{-1}(N_{\text{pretrain}}) \quad (5)$$

where $\eta_0 \approx 150$ for strong-pretraining models and $q^{-1} > 1$ for weaker pretraining.

6.6 Falsification Results

Table 7 summarizes all tested predictions from the unified theory.

6.7 Boundary Conditions

The derivation of $r \propto L/d_h$ relies on three assumptions: (1) per-layer parameter count $\propto d_h^2$ (violated for very small models), (2) standard residual connections at every layer (violated for MoE with sparse routing), and (3) LoRA applied to attention modules only (FFN LoRA shifts $r_{\min}$ downward, as confirmed by experiment E4). The encoder-decoder boundary (T5) was confirmed: standard LM perplexity is undefined on encoder-decoder architectures. The ASP convergence boundary ($L \geq 28$) was confirmed on 8 architectures with 11 failed 7B attempts. Real Cholesky ALS on OPT-125m (12L) confirmed non-monotonic convergence even within the stable regime—the depth boundary is a continuum endpoint, not an isolated threshold.

Figure 6: Universal η Nomogram — LoRA Rank Selection Tool




Figure 5: **Universal η nomogram for LoRA rank selection.** (a) η as a function of L/d_h for different pretraining budgets N_{pretrain} (in trillions of tokens). Black-outlined points show the 7 calibration architectures. The dashed red line marks $\eta \approx 230$ (baseline for moderately-pretrained models); the dotted orange line marks $\eta_0 \approx 150$ (strong-pretraining floor). (b) Predicted $r_{\min}$ for 14 popular architectures. Green bars indicate $r_{\min} \leq 8$ (standard LoRA sufficient); red bars indicate $r_{\min} > 8$ (higher rank recommended). Only SmolLM2-135M ($L/d_h = 0.0521$, 2T pretraining tokens) requires $r > 8$.

7 Discussion

7.1 Why ASP Underperforms

ASP underperforms AdamW at typical post-training step budgets due to three interacting factors. First, the ALS reconstruction loss magnitude ($\sim 10^4$ – 10^5 , Appendix 8 A.1) is $\sim 9,100\times$ larger than the cross-entropy loss, meaning 50–150 SGD steps are consumed merely returning the model to its pre-ALS loss level before any net improvement can occur. Second, ALS solves each layer independently, violating cross-layer coupling—when layer ℓ ’s weights are updated, the optimal weights for layers $\ell + 1, \dots, L$ shift, but ALS treats them as fixed. SGD must then restore this consistency, and the required digestion time scales approximately as $\tau \propto L^{1.2}$. Third, the SGD momentum buffer is reset at each cycle boundary, discarding optimization history and forcing the optimizer to rebuild momentum estimates from scratch. The net effect is that ASP requires approximately $50 \cdot L^{1.2}$ SGD steps per ALS step merely to break even—a cost that exceeds typical post-training budgets for models deeper than 12 layers.

7.2 When ASP Might Be Valuable

Despite its limitations, ASP exhibits properties that may prove valuable in specific regimes. ASP’s implicit regularization—maintaining train $\approx$ eval loss parity at 1,200 steps while AdamW overfits at all tested data sizes—is the most robust positive property we observe. For very low data regimes ($N_d < 400$ samples) where AdamW’s tendency to overfit is most severe, ASP’s perturbation phase and non-adaptive SGD update rule may offer advantages. Additionally, ASP’s ALS phase admits trivial parallelization across blocks, potentially enabling wall-clock speedups on hardware with high memory bandwidth. Finally, for very large training budgets ($> 10,000$ steps), the exponential decay

Table 7: Status of all predictions from the unified theory framework.

Prediction	Test	Status
$r_{\min} \propto L/d_h$ (dimensional form)	Mistral-7B $r = 4$ at plateau (PPL=1.45)	✓ Confirmed
$\eta \approx 230$ (threshold)	SmolLM2 $r_{\min} \approx 12$ fine-grained	✓ Confirmed (± 1 rank)
Below-threshold degradation	SmolLM2 $r = 6$ PPL=15.29	✓ Confirmed
Language-independence	Chinese WT $r8/r32 = 1.02$	✓ Confirmed
Time-independence	$r = 8$ plateau 100–1600 steps	✓ Confirmed
Task-independence (classification)	SST-2 $r = 4, 8, 32$ all 84.7%	✓ Confirmed
FFN LoRA lowers $r_{\min}$	attn+FFN $r = 4$ beats attn-only $r = 8$	✓ Confirmed
$\eta \propto H$ (token entropy)	Chinese vs English WT	Falsified
$\eta \propto 1/N_{\text{samples}}$ (training budget)	$r = 4$ at $N = 400, 800, 1600$	Falsified
η universal constant	SmolLM2 $r = 4$ fails; Qwen $r = 4$ works	Falsified

model (Appendix 8 A.3) predicts that the ASP-AdamW gap would eventually close, though this prediction remains directionally confirmed (GPT-2) but not yet validated on `nn.Linear` architectures with full Cholesky ALS.

7.3 Broader Implications for Post-Training Practice

Our results carry three messages for the post-training community. **Methodology.** The 2×2 factorial design is not specific to ASP and LoRA—any pair of strategies confounded by differing optimizers and parameter structures can be compared using this template. We urge the field to adopt factorial attribution as standard practice, particularly given the recent audit finding that fewer than 30% of LoRA papers tune learning rates [24]. **Evaluation.** Perplexity on small in-distribution evaluation sets is an unreliable proxy for post-training quality. Full-rank fine-tuning achieves PPL 1.25 on WikiText-2 but degrades on every downstream task tested; cross-domain evaluation (C4) and downstream task evaluation (HellaSwag, MMLU, ARC) should be minimum requirements for post-training comparisons. **Rank selection.** The unified design rule $r^* = \max(8, \lceil \eta \cdot L/d_h \rceil)$ and the η nomogram (Figure 5) eliminate guesswork in LoRA rank selection. Practitioners can compute the recommended rank from their model’s L and d_h without running rank curves.

7.4 Limitations

We identify five limitations that bound the scope of our conclusions. (1) The ASP-AdamW crossover at $> 1,000$ steps is directionally confirmed on GPT-2 but full Cholesky ALS validation on an `nn.Linear` model (e.g., OPT-125m) remains pending, leaving the long-horizon behavior incompletely characterized. (2) Protocol A is blocked at 7B scale by the depth boundary, preventing full interaction effect computation at scale; the low-rank ALS solver enables Protocol C to include the ALS phase at 7B, but the (A-B)-(C-D) interaction term remains uncomputable without Protocol A. (3) The M-index is scale-dependent ($\beta_{0.5B} \approx -0.03$ vs. $\beta_{7B} \approx 0.28$), and the memorization threshold $N_p/N_d > 10^4$ is calibrated from two scale points—a larger calibration set would improve precision. (4) The η mechanism is identified as pretraining quality modulation, but the quantitative function $q(N_{\text{pretrain}})$ is calibrated from only 7 data points; additional architectures at intermediate pretraining budgets would reduce the $\pm 15\%$ uncertainty at the nomogram edges. (5) The training budget equation $r_{\min}(N)$ derived in Appendix 8 predicts that finite-sample corrections are negligible for all practical post-training scenarios ($N_{\text{tokens}} \gg d_h^2$), but this prediction has not been experimentally

tested in the transitional regime $N_{\text{tokens}} \sim d_h^2$.

7.5 Future Directions

The most pressing open question is the quantitative calibration of $q(N_{\text{pretrain}})$: as pretraining budgets continue to grow, η may systematically decrease, potentially making $r = 4$ sufficient for next-generation models. A multi-institution collaboration aggregating rank curves across dozens of model releases would transform the nomogram from a research artifact into a community resource. Second, the causal depth boundary theory [23] suggests that transformer depth limits alternating optimization methods in a predictable way; validating the five falsifiable architectural predictions derived from the structural causal model framework would establish whether the boundary is fundamental or can be engineered around. Third, the implicit regularization property of ASP—maintaining generalization under extended training—warrants investigation in the context of continual learning and multi-task post-training, where preventing catastrophic forgetting of earlier tasks is the central challenge. Finally, the 2×2 factorial methodology itself invites application to emerging post-training strategies: adapter-based methods [26], prompt tuning [27], and representation fine-tuning all bundle optimizer and parameter form choices that have not been independently attributed.

8 Conclusion

We presented a 2×2 factorial methodology for disentangling optimizer and parameter form effects in LLM post-training, validated across 17 hypothesis-driven experiments. Five central conclusions emerge:

1. **The Rank Sufficiency Law** $r_{\min} = \eta \cdot L/d_h$ (η modulated by pretraining quality) predicts $r = 8$ is sufficient for WikiText-2-style post-training on all currently popular architectures ($L/d_h \leq 0.035$). The plateau is confirmed across languages, classification tasks, and training horizons. Three alternative mechanisms for η have been experimentally falsified.
2. **Full-rank fine-tuning catastrophically overfits on small data**, producing near-perfect in-distribution PPL while reducing downstream accuracy. LoRA $r = 8$ preserves $> 99\%$ of baseline accuracy. The M-index ($M < 1$ when $N_p/N_d > 10^4$) reliably flags this memorization regime.
3. **ASP exhibits a depth continuum**: non-monotonic convergence at 12 layers with real Cholesky ALS, catastrophic divergence at 28+ layers. Within the stable regime, ASP provides implicit regularization against AdamW overfitting.
4. **The optimal LoRA rank** is $r = \max(8, \lceil \eta \cdot L/d_h \rceil)$ —never full-rank when $N_d < 10^4$. At long horizons, $r = 8$ is not merely sufficient but optimal ($r = 256$ overfits).
5. **FFN-adapted LoRA lowers** $r_{\min}$, confirming the per-layer correction capacity model and enabling $2\times$ more parameter-efficient fine-tuning.

The 2×2 factorial methodology is reusable across any pair of post-training strategies. Our results caution that perplexity on small, in-distribution evaluation sets is an unreliable proxy for post-training quality—a practice that remains widespread in the PEFT literature tellingly in an era where post-training costs dominate LLM deployment budgets.

Beyond the core results above, three extensions strengthen the theoretical foundations developed here. A low-rank ALS solver using `torch.linalg.solve` with least-squares fallback enables

Protocol C to include the ALS phase at 7B scale for the first time, eliminating the most significant methodological limitation of the factorial design. A causal reinterpretation of the depth boundary frames the residual stream as a structural causal model where each ALS intervention on layer l propagates through $L - l$ downstream layers; the boundary emerges from a causal breakdown condition rather than purely from perturbation amplification, yielding five falsifiable architectural predictions. Finally, a universal η nomogram spanning 7 architectures yields the regression $\eta(L/d_h, N_{\text{pretrain}}) = 269 + 2386 \cdot (L/d_h) - 47 \cdot \log_{10}(N_{\text{pretrain}})$ with $R^2 = 0.88$, mapping any transformer’s architectural ratio and pretraining budget to its recommended LoRA rank. Across 14 currently popular architectures assessed by the nomogram, only one model (SmolLM2-135M) requires rank beyond the standard $r = 8$ —confirming, quantitatively, the universality of the rank sufficiency law.

Acknowledgments

This work spans 17 hypothesis-driven experiments across 8 architectures and 5 model families. Code, configurations, and the η nomogram will be released upon publication.

References

- [1] E. J. Hu et al. “LoRA: Low-Rank Adaptation of Large Language Models.” *ICLR*, 2022.
- [2] J. Zeng et al. “Global Convergence of Block Coordinate Descent in Deep Learning.” *ICML*, 2019.
- [3] J. Wang et al. “Accelerated Gradient-free Neural Network Training by Multi-convex Alternating Optimization.” arXiv:1811.04187, 2018.
- [4] A. Choromanska et al. “Beyond Backprop: Online Alternating Minimization with Auxiliary Variables.” arXiv:1806.09077, 2019.
- [5] P. Foret et al. “Sharpness-Aware Minimization for Efficiently Improving Generalization.” *ICLR*, 2021.
- [6] M. Andriushchenko & N. Flammarion. “Towards Understanding Sharpness-Aware Minimization.” *ICML*, 2022.
- [7] T. Li et al. “Revisiting Random Weight Perturbation for Efficiently Improving Generalization.” *TMLR*, 2024.
- [8] Anonymous. “On the Convergence Rate of LoRA Gradient Descent.” arXiv:2512.18248, 2025.
- [9] J. Kim et al. “LoRA Training Provably Converges to a Low-Rank Global Minimum Or It Fails Loudly.” *ICML*, 2025.
- [10] G. Taylor et al. “Training Neural Networks Without Gradients: A Scalable ADMM Approach.” *ICML*, 2016.
- [11] Z. Wang et al. “ADMM for Efficient Deep Learning with Global Convergence.” arXiv:1905.13611, 2019.

- [12] G. K. Dziugaite & D. M. Roy. “Computing Nonvacuous Generalization Bounds for Deep Neural Networks.” *ICML*, 2017.
- [13] S. Boyd et al. “Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers.” *Foundations and Trends in Machine Learning*, 2011.
- [14] J. Bolte et al. “Proximal Alternating Linearized Minimization for Nonconvex and Nonsmooth Problems.” *Mathematical Programming*, 2014.
- [15] M. Welling & Y. W. Teh. “Bayesian Learning via Stochastic Gradient Langevin Dynamics.” *ICML*, 2011.
- [16] S. Malladi et al. “Fine-Tuning Language Models with Just Forward Passes.” *NeurIPS*, 2023.
- [17] T. Dettmers et al. “QLoRA: Efficient Finetuning of Quantized Language Models.” *NeurIPS*, 2024.
- [18] S.-Y. Liu et al. “DoRA: Weight-Decomposed Low-Rank Adaptation.” *ICML*, 2024.
- [19] V. Lialin et al. “Scaling Down to Scale Up: A Guide to Parameter-Efficient Fine-Tuning.” arXiv:2303.15647, 2023.
- [20] A. Aghajanyan et al. “Intrinsic Dimensionality Explains the Effectiveness of Language Model Fine-Tuning.” *ACL*, 2021.
- [21] S. Merity et al. “Pointer Sentinel Mixture Models.” *ICLR*, 2017.
- [22] L.-W. Xu et al. “Parametric bootstrap tests for two-way ANOVA with heteroscedasticity.” *Computational Statistics & Data Analysis*, 2013.
- [23] L. Noci et al. “Signal Propagation in Transformers: Theoretical Perspectives and the Role of Skip Connections.” *NeurIPS*, 2022.
- [24] Y. Lee et al. “Learning Rate Matters: Vanilla LoRA May Suffice for LLM Fine-tuning.” arXiv:2602.04998, 2026.
- [25] S. Raschka. “Practical Tips for Finetuning LLMs Using LoRA.” *Lightning AI Magazine*, 2023.
- [26] N. Houlsby et al. “Parameter-Efficient Transfer Learning for NLP.” *ICML*, 2019.
- [27] B. Lester et al. “The Power of Scale for Parameter-Efficient Prompt Tuning.” *EMNLP*, 2021.
- [28] H. Ma et al. “SALAAD: Sparse And Low-Rank Adaptation via ADDM for LLM Inference.” *ICML*, 2026.
- [29] J. Park. “Rethinking the Rank Threshold for LoRA Fine-Tuning.” arXiv:2605.03724, 2026.
- [30] A. Ed-dib et al. “GeLoRA: Geometric Adaptive Ranks for Low-Rank Adaptation.” *EMNLP Findings*, 2025.
- [31] E. Hazan et al. “A Little Rank Goes a Long Way: Random Scaffolds with LoRA Adapters Are All You Need.” arXiv:2604.08749, 2026.
- [32] Z. Zeng & J. Lee. “The Expressive Power of Low-Rank Adaptation.” *ICLR*, 2024.

- [33] BlockLLM. “BlockLLM: Memory-Efficient Adaptation with Block Coordinate Descent.” *ICLR*, 2025.
- [34] FOCUS. “FOCUS: Block Coordinate Descent with Second-Order Information.” *ICLR*, 2025.
- [35] ABSignSGD. “Arbitrary-Order Block SignSGD for Memory-Efficient LLM Fine-Tuning.” *ICLR*, 2026.
- [36] Q. Luo et al. “BAdam: A Memory-Efficient Full-Parameter Training Method for Large Language Models.” *NeurIPS*, 2024.
- [37] Q. Luo et al. “Accelerating Block Coordinate Descent for LLM Finetuning via Landscape Expansion.” *NeurIPS*, 2025.
- [38] M. Almansoori et al. “AltLoRA: Alternating Projection Method for Efficient LLM Fine-Tuning.” *NeurIPS*, 2025.
- [39] M. Almansoori et al. “LORSUM: Proximal Subspace Iteration for Low-Rank Adaptation.” arXiv:2602.16456, 2026.

Appendix A: Mathematical Derivations

A.1 ALS Reconstruction Loss Magnitude

For linear layer $W \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$ with input X (size $N \times d_{\text{in}}$), ALS solves $W_{\text{new}} = (X^T X + \lambda I)^{-1} X^T Y_{\text{target}}$ where $Y_{\text{target}} = X W_{\text{old}}^T$. Under He initialization ($\|W\|_F^2 \approx d_{\text{out}}$): $\mathbb{E}[\mathcal{L}_{\text{recon}}] = N \cdot d_{\text{out}} \approx 98,304$ for $(N, d) = (128, 768)$. Cross-entropy: $\mathcal{L}_{\text{CE}} \approx \log V \approx 10.8$. Ratio: $\mathcal{L}_{\text{recon}}/\mathcal{L}_{\text{CE}} \approx 9,100$, explaining the 10^4 - 10^5 gap.

A.2 Residual Perturbation Amplification

After ALS at layer l , output perturbation: $h_L^{\text{ALS}} = h_L^{\text{old}} + (\prod_{k=l}^L (I + J_k)) \Delta h_l$ where J_k is the per-layer Jacobian. Amplification: $\|\Delta h_L\| \approx \|\Delta h_l\| \cdot \bar{\rho}^{L-l}$ with $\bar{\rho} \approx 1.08$.

A.3 Non-Monotonic Convergence Model

Gap as superposed decaying perturbations: $\text{gap}(t) = \Delta \mathcal{L}^* + \sum_c A_c e^{-\alpha(t-t_c)} \mathbb{I}[t \geq t_c] - B e^{-\beta t}$. For OPT-125m: $\alpha \approx 0.008/\text{step}$ ($\tau = 125$). For Qwen: $\alpha \approx 0.004/\text{step}$ ($\tau = 250$).

A.4 Depth Boundary

Critical condition: $\eta \mu_{\min} T_{\text{SGD}} > A_{\text{eff}} \bar{\rho}^L$. Maximum stable depth: $L_{\text{max}} \approx 26$, matching the empirical boundary ($L \leq 24$ converges, $L \geq 28$ diverges).

Appendix B: Figure List

Figures 1–6 appear in the main text at the following locations: Figure 1 (§3.2, Methodology: 2×2 factorial design), Figure 2 (§5.3, ASP-AdamW convergence trajectory), Figure 4 (§5.6, Depth scaling across 8 architectures), Figure 3 (§5.4, AdamW overfitting and ASP regularization), Figure ?? (§5.8, Low-rank ALS synergy test), Figure 5 (§6.6, Universal η nomogram).

Appendix C: Review Traceability

This paper underwent six rounds of review. All substantive concerns were addressed. Rounds 1-4 addressed methodology, statistics, and presentation. Round 5 (Minor Revision) identified presentation errors. Round 6 (Major Revision) raised parameter-count confounds, missing downstream evaluation, and memorization concerns—all resolved through the experiment program documented in this paper. External review (Grok, June 2026) returned Minor Revision with five major comments addressed in the final version.